



Critical Sections: Implementing Logical Atomicity

- Critical section of code is one that:
 - must be executed by one thread at a time
 - otherwise, there is a race condition
 - Example: linked list from before
 - Insert/Delete code forms a critical section
 - What about just the Insert *or* Delete code?
 - is that enough, or do both procedures belong in a single critical section?


List Example




Insert: $\text{elem} \rightarrow \text{next} := \text{head};$




Remove: $t := \text{head};$



Insert: $\text{head} := \text{elem};$



Remove: $\text{head} := \text{head} \rightarrow \text{next};$



Remove: return t ;

Critical Section (CS) Problem

- Provide entry and exit routines:
 - all threads must call entry before executing CS
 - all threads must call exit after executing CS
 - a thread must not leave entry routine until it's safe
- CS solution properties (**red**: safety; **black**: liveness)
 - **Mutual exclusion**: at most one thread is executing CS
 - **Absence of deadlock**: two or more threads trying to get into CS, and no thread is in \rightarrow at least one succeeds
 - **Absence of unnecessary delay**: if only one thread tries to get into CS, and no thread is in CS, it succeeds
 - **Eventual entry**: thread eventually gets into CS

Structure of threads for Critical Section problem

Threads do the following:

```
while (1) {  
    do “other stuff” (i.e., non-critical section)  
    call CS entry routine  
    execute CS  
    call CS exit routine  
}
```

Critical Section Assumptions

- Threads must call CS entry and exit routines
- Threads must not die or quit inside a critical section (or the entry and exit routines)
- Threads **can** be context switched inside a critical section
 - this does **not** mean that another thread may enter the critical section

Critical Section Solution Attempt #1

(2 thread version, with id's 0 and 1)

```
turn = 0
```

```
entry(id) {  
    while (turn != id) ;  
}
```

```
exit(id) {  
    turn := 1-id;  
}
```

Critical Section Solution Attempt #2

(2 thread version, with id's 0 and 1)

```
flag[0] = flag[1] = false
```

```
entry(id) {  
    flag[id] := true;  
    while (flag[1-id]) ;  
}
```

```
exit(id) {  
    flag[id] := false;  
}
```

Critical Section Solution Attempt #3

flag[0] = flag[1] = false, turn == 0

```
entry(id) {  
    flag[id] := true;  
    turn := 1-id;  
    while (flag[1-id] and turn == 1-id) ;  
}  
  
exit(id) {  
    flag[id] := false;  
}
```


Summary: Satisfying the 4 properties (on Attempt #3)

- Mutual exclusion
 - turn must be 0 or 1 \Rightarrow only one thread can be in CS
- Absence of deadlock
 - turn must be 0 or 1 \Rightarrow one thread will be allowed in
- Absence of unnecessary delay
 - only one thread trying to get into CS \Rightarrow flag[other] is false \Rightarrow the thread will get into the CS
- Eventual Entry
 - spinning thread will not modify turn
 - thread trying to go back in will set turn equal to spinning thread's id

Next three slides are just for
reference (include comments)

Critical Section Solution Attempt #1

(2 thread version, with id's 0 and 1)

```
turn = 0 /* turn is shared */  
  
entry(id) { /* note id local to each thread */  
    while (turn != id) ; /* if not my turn, spin */  
}  
  
exit(id) {  
    turn := 1-id; /* other thread's turn */  
}
```

Critical Section Solution Attempt #2

(2 thread version, with id's 0 and 1)

```
flag[0] = flag[1] = false
```

```
/* flag is a shared array */
```

```
entry(id) {
```

```
    flag[id] := true; /* I want to go in */
```

```
    while (flag[1-id]) ; /*proceed if other not trying*/
```

```
}
```

```
exit(id) {
```

```
    flag[id] := false; /* I'm out */
```

```
}
```

Critical Section Solution Attempt #3

(2 thread version, with id's 0 and 1)

flag[0] = flag[1] = false, turn == 0

/* flag and turn are shared variables */

entry(id) {

 flag[id] := true; /* I want to go in */

 turn := 1-id; /* in case other thread wants in */

 while (flag[1-id] and turn == 1-id) ;

}

exit(id) {

 flag[id] := false; /* I'm out */

}

Hardware Support for Critical Sections

- All modern machines provide atomic instructions
- All of these instructions are *Read/Modify/Write*
 - atomically read value from memory, modify it in some way, write it back to memory
- Use to develop simpler critical section solution **for any number of threads**

Example: Test-and-Set

Some machines have it

```
function TS(ref target: bool) returns bool  
  bool b := target; /* return old value */  
  target := true;  
  return b;
```

Executes atomically

CS solution with Test-and-Set

```
s = false /* s is a shared variable */
```

```
entry( ) {
```

```
    bool spin; /* spin is local to each thread! */
```

```
    spin := TS(s);
```

```
    while (spin)
```

```
        spin := TS(s);
```

```
}
```

```
exit( ) {
```

```
    s := false;
```

```
}
```

Function TS(bool &target) returns bool

```
    bool b := target
```

```
    target := true
```

```
    return b
```


A few example atomic instructions


- Compare and Swap (x86)
- Load linked and conditional store (RISC)
- Fetch and Add
- Atomic Swap
- Atomic Increment

Basic Idea with Atomic Instructions

- One shared variable plus per-thread flag
- Use atomic instruction with flag, shared variable
 - On a change, allow thread to go in
 - Other threads will not see this change
 - Should use “test and test and set” style (see next slide)
- When done with CS, set shared variable back to initial state
 - Some advanced solutions spin on exit (not discussed in 422)

CS using Test-and-Test-and-Set

Initially, $s == \text{false}$

```
entry( ) {  
    bool spin;  
    spin := TS(s);  
    while (spin) {  
        while (s) ;  This line is the difference from “regular” Test-and-Set  
        spin := TS(s);  
    }  
}  
  
exit( )  
    s := false;
```

Function TS(bool &target) returns bool
bool b := target
target := true
return b

The (Fair) Ticket Algorithm-- High Level Solution

Init: $\text{number} = \text{next} = 0$

Entry: $\langle \text{ticket} = \text{number}; \text{number}++ \rangle$
 $\langle \text{await} (\text{ticket} == \text{next}) \rangle$

Exit: $\langle \text{next}++ \rangle$

Fair Ticket Algorithm--Implementation

Function FA(ref int target, const int C) returns int
int loc := target
target := target + C
return loc

number = next = 0

entry() {

int ticket = FetchAndAdd(number, 1)

while (ticket != next) ;  No atomicity needed---why?

}

exit()

next = next + 1  No atomicity needed---why?

Works for any number of threads---and is fair (provides eventual entry)

N-Thread CS Solutions *Without* Atomic Instructions

- Much more complex
 - We will not address these, and you will not be responsible for them
- Most common solution is called the bakery algorithm
 - Essentially an implementation of the ticket algorithm, but without atomic instructions (other than load and store, of course)
 - If interested, see:
https://en.wikipedia.org/wiki/Lamport%27s_bakery_algorithm

Problems with busy-waiting CS solution

- Complicated
- Inefficient
 - consumes CPU cycles while spinning
- Priority inversion problem
 - low priority thread in CS, high priority thread spinning can end up causing deadlock
 - example: Mars Pathfinder problem

In some cases, want to block when waiting for CS

Locks

- Two operations:
 - Acquire (get it, if can't go to sleep or spin)
 - Release (give it up, possibly wake up a waiter or spinning thread)
- Acquire and Release are (logically) atomic
- **A thread can only release a previously acquired lock**
- `entry()` is then just `Acquire(lock)`
- `exit()` is just `Release(lock)`

Lock is shared among all threads

Using Locks for Critical Sections (CSs)

Assume L1 and L2 are locks; x and y are shared variables

- Single CS, one location
 - Acquire(L1)
 - CS (e.g., $x += k$)
 - Release(L1)
- Single CS, two locations
 - Acquire(L1)
 - CS (e.g., $x += k_1$)
 - Release(L1)
 - ...
 - Acquire(L2)
 - CS (e.g., $x += k_2$)
 - Release(L2)
- Two different CSs
 - Acquire(L1)
 - CS 1 (e.g., $x += k_1$)
 - Release(L1)
 - ...
 - Acquire(L2)
 - CS 2 (e.g., $y += k_2$)
 - Release(L2)

Incorrect Blocking Lock Implementation in an OS kernel (single-core machine)

- Acquire(lock) disables interrupts
- Release(lock) enables interrupts
- Advantages:
 - is a blocking solution; can be used inside OS in some situations
- Disadvantages:
 - CS can be in user code [could infinite loop], might need to access disk in middle of CS, system clock could be more skewed than typical, etc.

Correct Blocking Lock Implementation in an OS kernel (single-core machine)

lock class has queue, value

Initially:

queue is empty

value is free

Acquire(lock)

Disable interrupts

if (lock.value == busy)

 enQ(lock.queue, thread)

 go to sleep

else

 lock.value := busy

Enable interrupts

Release(lock)

Disable interrupts

if notEmpty(lock.queue)

 thread := deQ(lock.queue)

 enQ(readyList, thread)

else

 lock.value := free

Enable interrupts

Can interrupts be enabled before sleep?

lock class has queue, value

Initially:

queue is empty

value is free

Acquire(lock)

Disable interrupts

if (lock.value == busy)

Enable interrupts

enQ(lock.queue, thread)

go to sleep

else

lock.value := busy

Enable interrupts

Release(lock)

Disable interrupts

if notEmpty(lock.queue)

thread := deQ(lock.queue)

enQ(readyList, thread)

else

lock.value := free

Enable interrupts

Can interrupts be enabled before sleep?

lock class has queue, value

Initially:

queue is empty

value is free

Acquire(lock)

Disable interrupts

if (lock.value == busy)

enQ(lock.queue, thread)

Enable interrupts

go to sleep

else

lock.value := busy

Enable interrupts

Release(lock)

Disable interrupts

if notEmpty(lock.queue)

thread := deQ(lock.queue)

enQ(readyList, thread)

else

lock.value := free

Enable interrupts

What about a “spin-lock”?

Items in red must be addressed

lock class has queue, value

Initially:

queue is empty

value is free

Acquire(lock)

Disable interrupts

if (lock.value == busy)

 enQ(lock.queue, thread)

 go to sleep

else

 lock.value := busy

Enable interrupts

Release(lock)

Disable interrupts

if notEmpty(lock.queue)

 thread := deQ(lock.queue)

 enQ(readyList, thread)

else

 lock.value := free

Enable interrupts

Fair Ticket *Lock*, Revisited

```
number = next = 0
```

```
Acquire( ) {
```

```
    int ticket = FetchAndAdd(number, 1)
```

```
    while (ticket != next) ;
```

```
}
```

```
Release( ) {
```

```
    next = next + 1
```

```
}
```

This is known as a “spin-lock”.

What about *blocking*, OS-level implementation on multiple cores?

- More complicated
 - Cannot use disabling/enabling interrupts
- Can be solved by using atomic instructions and spin locks together
 - If interested, consult the internet and/or AI tools

What about *blocking*, **user-level** implementation on multiple cores?

- Typically solved with atomic instructions and futex (**F**ast **U**erspace **muT**EX) support
 - Atomic instruction used for a thread wanting to acquire the lock
 - Futex provides wait and wake operations that use the OS to block the thread if lock is held by another thread
 - All the complication is pushed into the OS
- Again, consult the internet and/or AI tools for details

Blocking Locks vs Spin Locks

- Blocking Locks
 - Advantages: do not consume CPU cycles when spinning
 - Disadvantage: context switches to block and unblock
- Spin Locks
 - Advantages: faster when multiple cores and no competing jobs
 - Disadvantage: consumes CPU cycles when spinning; priority inversion possible

Problem with Locks

- Not general
 - Locks only solve critical section problem
 - Locks do not allow any more general synchronization
 - Specifically, locks cannot enforce strict orderings between threads

Condition synchronization

- Need to wait until some condition is true
- Example: thread join
- Example: bounded buffer (later)

Barriers

- Points in program at which all threads must arrive before any can proceed
- Barriers provide *sequence control*
- In CSC 422, we will assume that all threads must participate in a barrier
 - It is possible to have “local” barriers, in which a subset of threads participate
 - In the message passing paradigm, we call these subcommunicators (will not discuss in this class)

Challenges With Barriers

- Must implement barriers
- Must insert barriers into programs
 - Want as few barriers as possible
 - Barriers are overhead
 - However, must *ensure* correct program behavior

Centralized Barrier

Shared variable: $\text{count} = 0$; assume n threads

High-level code for barrier for a given thread:

```
<count++>
```

```
<await (count == n) ;>
```

Actual implementation for barrier for a given thread:

```
FetchAndAdd(count, 1)
```

```
while (count != n) ;
```

Problem?

Centralized Barrier

Shared variable: $\text{count} = 0$; assume n threads

High-level code for barrier for a given thread:

```
<count++>
```

```
<await (count == n) ;>
```

Actual implementation for barrier for a given thread:

```
FetchAndAdd(count, 1)
```

```
while (count != n) ;
```

Problem: **reset**

Centralized Barrier, Suggested by Past Students

Shared variable: `count = 0`

Code for barrier for a given thread:

```
FetchAndAdd(count, 1)
```

```
while (count mod n != 0) ;
```

Problem?

Centralized Barrier, Suggested by Past Students

Shared variable: `count = 0`

Code for barrier for a given thread:

```
FetchAndAdd(count, 1)  
while (count mod n != 0) ;
```

Problem: one or more threads might not see
barrier exit condition (`count mod n == 0`)

Centralized Barrier, Suggested by Past Students, Version 2.0

Shared variable: `count = 0`

Code for barrier for a given thread:

```
if (FetchAndAdd(count, 1) == n-1)
```

```
    count = 0
```

```
else
```

```
    while (count != 0) ;
```

FetchAndAdd returns the old value



Problem?

Centralized Barrier, Suggested by Past Students, Version 2.0

Shared variable: `count = 0`

Code for barrier for a given thread:

```
if (FetchAndAdd(count, 1) == n-1)
```

```
    count = 0
```

```
else
```

```
    while (count != 0) ;
```

FetchAndAdd returns the old value



Problem: same as previous attempt: a thread may not see the barrier exit condition

Centralized Barrier, Suggested by Current Students

Shared variable: `count = 0`

```
Barrier(ref int k) {    // separate k for each thread
    FetchAndAdd(count, 1)
    while (count < k * n);
    k = k + 1
}
```

Thread invokes via `Barrier(barrierNum)`, where `barrierNum` is kept by each thread and initially 0

This works, but has disadvantage that each thread must keep track of barrier state

Centralized Barrier (handles reset)

Shared variables: $\text{countEven} = \text{countOdd} = \text{nB} = 0$

Code for barrier for a given thread (assume n threads):

```
if (nB mod 2 == 0) {  
    if (FetchAndAdd(countEven, 1) == n-1) {  
        nB = nB + 1  
        countEven = 0  
    }  
    else  
        while (countEven != 0) ;  
    else {  
        // same code, but with countOdd instead of countEven  
    }  
}
```

Barrier from
centralized
version 2.0

FetchAndAdd returns the old value

Centralized Barrier (handles reset)

Full code

Shared variables: `countEven = countOdd = nB = 0`

Code for barrier for a given thread (assume n threads):

```
if (nB mod 2 == 0) {  
    if (FetchAndAdd(countEven, 1) == n-1) {  
        nB = nB + 1  
        countEven = 0  
    }  
    else  
        while (countEven != 0) ;  
    else {  
        if (FetchAndAdd(countOdd, 1) == n-1) {  
            nB = nB + 1  
            countOdd = 0  
        }  
        else  
            while (countOdd != 0) ;  
    }  
}
```



FetchAndAdd returns the old value

Centralized Barrier (handles reset)

Sense-reversing version

Shared variables: $\text{count} = 0$, $\text{globalSense} = \text{false}$

Code for barrier for a given thread (assume n threads):

```
bool localSense = not globalSense
```

```
if (FetchAndAdd(count, 1) == n-1) {
```

```
    count = 0
```

```
    globalSense = localSense
```

```
}
```

```
else
```

```
    while (globalSense != localSense) ;
```

FetchAndAdd returns the old value



Centralized Barrier (with reset) works, but...

- May not perform well
 - If we are implementing a barrier, we should assume it will be used frequently and possibly with a large number of threads
 - Spinning on a shared flag (*count*) is not ideal on some multicore architectures
 - Typically ones with higher core counts
 - We want to investigate barriers where each thread spins on a *different* variable

Symmetric Barrier, 2 threads (not quite correct)


```
int arrive[2] = {0,0}
```

Thread 0's code

```
arrive[0] = 1  
while (arrive[1] != 1)  
;  
arrive[1] = 0
```

Thread 1's code

```
arrive[1] = 1  
while (arrive[0] != 1)  
;  
arrive[0] = 0
```



Flag synchronization principle (1): a thread resets the flag on which it spins.

Symmetric Barrier, 2 threads (not quite correct)

```
int arrive[2] = {0,0}
```

Thread 0's code

```
arrive[0] = 1
```

```
while (arrive[1] != 1)
```

```
;
```

```
arrive[1] = 0
```

Thread 1's code

```
arrive[1] = 1
```

```
while (arrive[0] != 1)
```

```
;
```

```
arrive[0] = 0
```

Problem!



Flag synchronization principle (1): a thread resets the flag on which it spins.

Flag synchronization principle (2): a thread should not set a flag until it's known to be clear.

Symmetric Barrier, 2 threads (correct)

```
int arrive[2] = {0,0}
```

Thread 0's code

```
while (arrive[0] != 0)
;
arrive[0] = 1
while (arrive[1] != 1)
;
arrive[1] = 0
```

Thread 1's code

Implements flag synchronization principle (2)

```
while (arrive[1] != 0)
;
arrive[1] = 1
while (arrive[0] != 1)
;
arrive[0] = 0
```

Symmetric Barrier, 2^p threads

- Conceptually, just glue multiple two-thread barriers together

Symmetric Barrier, 2^p threads

- Conceptually, just glue multiple two-thread barriers together
 - Problem: flags meant for one thread might be seen by another thread
 - Two solutions to this:
 - use more storage (e.g., a two-dimensional array of arrive flags, with the first index as the round), or
 - set the arrive array values to the round id (simpler and more space efficient)

Dissemination Barrier

(Hensgen, Finkel, and Manber)

- Similar to symmetric barrier in that it has a logarithmic number of steps (in the number of participating threads)
- Different from symmetric barrier in that at each round, a thread will signal one thread and wait for a *different* thread's signal

Dissemination Barrier

`int arrive[0:P-1] = {0, 0, ..., 0}` in C, need *volatile*

Thread i 's code (local vars: j and $waitFor$):

```
for  $j = 1$  to ceiling( $\log_2(P)$ ) {  
    while ( $arrive[i] \neq 0$ ) ;  
     $arrive[i] = j$   
     $waitFor = (i + 2^{j-1}) \bmod P$   
    while ( $arrive[waitFor] \neq j$ ) ;  
     $arrive[waitFor] = 0$   
}
```