

# So far in this class...

- We have assumed a single address space
  - Threads have communicated via shared variables
    - E.g., bounded buffer and readers/writers problems
  - Synchronization objects (locks, barriers, semaphores, monitors) have been placed in shared memory
    - Accessible by all threads
- What if we want threads in two different address spaces to communicate?
  - Cannot simply use a synchronization object to communicate, because they do **not** share memory

# Message Passing

- Thread  $T_1$  in address space  $A_1$  sends a message to thread  $T_2$  in address space  $A_2$ 
  - Where are these threads/address spaces located?
    - Same machine? Different machine? Does it matter?
    - $A_1$  and  $A_2$  certainly *could* be on different machines; if so, the message must travel over a network; how does it get there?
    - Does the programmer have to worry about this?
  - What is in the message?
- Because a thread plus an address space is a process, sometimes we will refer to processes  $P_1$  and  $P_2$  in the above

# Message Passing

- In this class:
  - Message passing can occur only between threads in different address spaces
    - Again, we will use the term *processes* often
  - Message passing cannot occur between threads in the same address space. Why?

# Message Passing

- In this class:
  - Message passing can occur only between threads in different address spaces
    - Again, we will use the term *processes* often
  - Message passing cannot occur between threads in the same address space. Why?
    - Should use shared variables for communication between threads in the same address space
    - Could also be thought of as a process sending to and receiving from itself

# Physical Reality of Networks

- Networks are unreliable
  - messages are divided into packets
  - packets can get lost
  - packets can arrive out of order
  - receiver can get overloaded
    - cannot handle rate of packet arrival

# Define a new abstraction: Channels

- Analogous to abstractions in OS's
  - process -- abstraction of a processor
  - virtual memory -- abstraction of unlimited memory
  - files -- abstraction of disk
- Want to abstract communication network
  - don't want to worry about lost packets, out-of-order packet arrival, overflowing buffers, etc.
  - **Channel** -- abstraction of point to point, reliable communication link

# Send and Receive

- `Send(channel, expr1, expr2, ..., exprN)`
  - Sends a message on the channel indicated
  - Expressions *expr1*, *expr2*, ..., *exprN* can be l-vals or r-vals
- `Receive(channel, arg1, arg2, ..., argN)`
  - Receive a message from the channel indicated into *arg1*, *arg2*, ..., *argN*
  - Variables *arg1*, *arg2*, ..., *argN* must be l-vals (must provide a storage location)
  - Receive blocks until data is available on channel
    - Can be relaxed in real implementations

# “Ping-Pong” example

## chan PingPong[2](int x) is shared

Process 0 code

```
int a = 10, b
```

```
Send(PingPong[1], a)
```

```
Receive(PingPong[0], b)
```

```
print a, b
```

Process 1 code

```
int a
```

```
Receive(PingPong[1], a)
```

```
Send (PingPong[0], a)
```

```
print a
```

Notes:

- Process 0 sends to Process 1; Process 1 receives and sends back
- Output is 10, 10 for Process 0; 10 for Process 1
- Need two channels to avoid self-send/self-receive
- Note that Receive takes its variable as an l-val



# Ping-Pong code, shorthand

## Will be the format on quizzes

Process 0 code

```
int a = 10, b
```

```
Send(1, a)
```

```
Receive(1, b)
```

```
print a, b
```

Process 1 code

```
int a
```

```
Receive(0, a)
```

```
Send (0, a)
```

```
print a
```

Notes:

- Exact same idea as the version with channels
- Difference is Send names the process to which message is sent, and Receive names the process from which message is received
- Send/Receive match if each names the other and number and types of arguments match

# Exchange example

## chan Exchange[2](int x) is shared

Process 0 code

```
int a = 10, b
```

```
Send(Exchange[1], a)
```

```
Receive(Exchange[0], b)
```

```
print a, b
```

Process 1 code

```
int a, b = 8
```

```
Send(Exchange[0], b)
```

```
Receive(Exchange[1], a)
```

```
print a, b
```

Notes:

- Both processes send to the other and then receive from the other
- Output is 10, 8 for both processes
- Need two channels to avoid self-send/self-receive
- Note that Receive takes its variable as an l-val

# Exchange code, shorthand

## Will be the format on quizzes

Process 0 code

```
int a = 10, b
```

```
Send(1, a)
```

```
Receive(1, b)
```

```
print a, b
```

Process 1 code

```
int a, b = 8
```

```
Send(0, b)
```

```
Receive(0, a)
```

```
print a, b
```

Notes:

- Exact same idea as the version with channels
- Difference is Send names the process to which message is sent, and Receive names the process from which message is received

# Example Quiz Question

Assume **each process** has two integers, a and b, initialized to 0

What are the final values of a and b in each process?

Does the code terminate?

P0 code

Send(1, a+1)

Receive(1, b)

P1 code

Receive(0, b)

Send(2, b+1)

Receive(2, a)

Send(0, a+1)

P2 code

Receive(1, b)

Send(1, b+1)

# Another example

## chan S(int x) is shared

Process 0 code

```
Send(S, 10)
```

Process 1 code

```
int a = 5; Receive(S, a); print "I am P1", a
```

Process 2 code

```
int a = 5; Receive(S, a); print "I am P2", a
```

This will deadlock. What happens if there is another send in P0, of a different value? Which process receives which value?

# Send and Receive

- Notes:
  - Channel handles reliability
    - Must be implemented by network protocols
  - Access to channel is atomic
  - Message has to be buffered if it arrives but receiver has not yet invoked receive
    - In fact, receiver's view of channel is a FIFO queue of pending messages
  - Implementation requires synchronization
  - Send, Receive can be OS kernel primitives or can be library primitives (e.g., MPI, the library we will use for program 3)

# Send and Receive

- Notes, continued:
  - Special case: both *exprs* and *args* are the empty set; in this case:
    - Send is analogous to  $V(s)$
    - Receive is analogous to  $P(s)$
    - Number of pending “messages” is analogous to the value of  $s$ 
      - More precisely, number of pending invocations

# Duality of Monitors/Message Passing

- First observed by Lauer and Needham in 1978
  - “On the Duality of Operating System Structures”
  - Observed that a monitor program can be translated mechanically into a message passing program (and vice versa)



# Resource Allocation with Monitors

```
monitor ResourceAllocator
```

```
  int free = true; cond c
```

```
  acquire( ):    if (free) free = false  
                  else wait(c)
```

```
  release( ):    if (empty(c)) free = true  
                  else signal(c)
```

```
end ResourceAllocator
```

Client calls ResourceAllocator.acquire( )/ResourceAllocator.release( )

# Resource Allocation with Message Passing

## Server-side code

```
enum reqType := {ACQUIRE, RELEASE}
chan request(int clientId, reqType which)
chan reply[n]( ) // one entry per client, no params
Allocator ( )    // runs on server
    queue pending; // initially empty
    int clientId; bool free := true
    reqType which;
```

# Resource Allocation with Message Passing

## Server-side code, continued

```
while (1) {  
    receive(request, clientId, which)  
    switch(which) {  
        ACQUIRE:  
            if (free)  
                free := false; send(reply[clientId])  
            else  
                pending.insert(clientId)  
        RELEASE:  
            if notempty(pending) send(reply[pending.remove( )])  
            else free := true  
    }  
}
```

# Resource Allocation with Message Passing

## Client-side code

```
Client (i) {  
    send request (i, ACQUIRE)  
    receive reply[i]( ) // blocks awaiting permission  
    send request (i, RELEASE) // no block here  
}
```

# Duality of Monitors/Message Passing

## Monitors

Monitor variables

Entry (implicit mutex)

Procedures in monitor

Procedure call

Procedure return

Wait

Signal

## Message Passing

Local vars on server

Blocking recv on server

Arms of switch stmt

Client sends request

to server; may block  
awaiting reply

Server sends result to

proper client if necessary

Insert request on server Q

Remove & process request  
from server queue

# Resource Allocation with Monitors

**monitor ResourceAllocator**

**int free = true; cond c**

**acquire( ):**      if (free) free = false

                 else wait(c)

**release( ):**      if (empty(c)) free = true

                 else signal(c)

**end ResourceAllocator**

Client calls **ResourceAllocator.acquire( )/ResourceAllocator.release( )**

# Resource Allocation with Message Passing

## Server-side code

```
enum reqType := {ACQUIRE, RELEASE}  
chan request(int clientId, reqType which)  
chan reply[n]( ) // one entry per client  
Allocator ( ) // runs on server  
    queue pending; // initially empty  
    int clientId; bool free := true  
    reqType which;
```

# Resource Allocation with Message Passing

## Server-side code, continued

```
while (1) {  
    receive(request, clientId, which)  
    switch(which) {  
        ACQUIRE:  
            if (free)  
                free := false; send(reply[clientId])  
            else  
                pending.insert(clientId)  
        RELEASE:  
            if notempty(pending) send(reply[pending.remove( )])  
            else free := true  
    }  
}
```



# Resource Allocation with Message Passing

## Client-side code

```
Client (i) {  
    send request (i, ACQUIRE)  
    receive reply[i]( ) // blocks  
    send request (i, RELEASE) // no block here  
}
```

# Programming Client/Server Applications (General Outline)

## Outline of Client code

```
while (1) {  
    build request  
    send(request, server)  
    receive(reply)  
    do something  
}
```

## Outline of Server code

```
while (1) {  
    receive(request)  
    switch(request)  
    case type1:  
        send(client, reply1)  
    case type2:  
        send(client, reply2)  
    etc.  
}
```