

# Problems with Send and Receive

- Low level
  - programmer is engaged in I/O
  - server often not modular
  - takes 2 calls to get what you want (send, followed by receive) -- error prone
- Solution
  - use procedure calls -- familiar model

# Remote Procedure Call (RPC)

(called RMI in Java)

- Allow procedure calls to other machines
  - servicing of procedure remote
  - caller blocks until procedure finished, as usual
  - simpler than explicit message passing
- Complications
  - caller and receiver in different address spaces
  - parameter passing
  - where is the machine providing the service?
  - what about crashes?

# Programming Client/Server Applications (General Outline)

## Outline of Client code

```
while (1) {  
    build request  
    send(request, server)  
    receive(reply)  
    do something  
}
```

## Outline of Server code

```
while (1) {  
    receive(request)  
    switch(request.type)  
    case FOO:  
        ...  
        send(client, reply1)  
    case BAR:  
        ...  
        send(client, reply2)  
    etc.
```

# Programming Client/Server Applications with RPC

## Outline of Client code

```
while (1) {  
    reply = foo(params)  
    do something  
}
```

Server is written as  
collection of several  
procedures

## Outline of Server code

```
foo(params) {  
    ....  
    return reply1;  
}  
  
bar (params) {  
    ....  
    return reply2;  
}
```

# Basics of RPC Implementation

- Goal: provide complete transparency to RPC user
  - RPC implementation must make a remote call look like a local (regular) procedure call
  - **Implementation** replaces a normal procedure call with:
    - pack arguments (including function to be executed) into a message via a “stub function”
      - may need to worry about dynamic data structures as well as low-level concerns such as byte ordering
    - send message to server; block waiting for reply
      - implemented via explicit message passing (send/receive)

# RPC User's View

Client

```
while (1) {  
    reply = foo(params)  
    do something  
}
```

Server

```
foo(params) {  
    ....  
    return reply1  
}
```

# RPC Implementation

Client

```
while (1) {  
reply = foo(params)  
  foo_stub(params, &reply)  
  do something  
}
```

```
foo_stub(params, &reply) {  
  msg.func = foo  
  msg.data[0] = param1  
  msg.data[1] = param2  
  send(Server, msg)  
  receive(Server, result)  
  reply = result.returnVal  
}
```

Server

```
foo(params) {  
  ....  
  return reply1  
}
```

User does not know  
anything about foo\_stub

User does not know  
anything about foo\_stub

# Basics of RPC Implementation

- Goal: provide complete transparency
  - On receipt at server: unpack and push parameters onto the stack, call function (**create new thread**)
    - Implemented by creating a thread that calls a stub function
  - Server then sends reply to client with results of function
  - On receipt of reply at client: put result where it belongs, unblock client



# RPC Implementation

## Client

```
while (1) {  
    reply = foo(params)  
    foo_stub(params, &reply)  
    do something  
}  
  
foo_stub(params, &reply) {  
    msg.func = FOO  
    msg.data[0] = param1  
    msg.data[1] = param2  
    send(Server, msg)  
    receive(Server, result)  
    reply = result.returnVal  
}
```

## Server

```
foo(params, &returnVal) {  
    ....  
    return reply1  
    returnVal = reply1  
}  
  
RPC_server( ) {  
    receive((Client = ANY_SOURCE), msg)  
    switch(msg.func) {  
        case FOO:        User does not know  
                           anything about RPC_server  
            t = thread_create(foo_stub', params)  
    }  
  
    foo_stub'(params) {  
        User does not know  
        anything about foo_stub'  
        foo(params, returnVal)  
        msg.returnVal = returnVal  
        send(Client, msg)  
    }  
}
```

# RPC Implementation Issues

- Weakly typed languages
  - E.g., C --- what to do if unbounded array passed to RPC?
  - Pointers across different machines?
- Communication via global variables impossible
- Binding
  - How does client know where servicing machine is?
    - One solution: use a database
- Failures?
  - What if function is partially executed, or executed twice, or executed never?

# RPC Parameter Passing

- Client machine may be a different architecture than server
  - we will ignore this issue – one side must convert data if byte ordering is an issue
- Parameter issues
  - what parameter passing style should be provided?
  - can be important performance issue
  - not as easy as it seems at first glance

# Call by Value

- Simple semantics
- Just package up the args, and send them
  - can be problematic (efficiency-wise) if pointer parameter points to a complex data type, e.g., graph or list
- Server uses these args
  - doesn't need to send them back

# Call by Reference

- What do pointers mean across machines?
  - they mean nothing across address spaces on the same machine, let alone on different machines
- Could send back message to client on each reference
  - **Would be slow**
  - Never used for RPC

# Call by Copy/Restore

- Similar to call by reference
  - parameter copied in, same as call by value
    - same disadvantages of having to copy entire structures
  - but when procedure finished, copy parameter back to caller
  - not quite the same as call by reference
  - method of choice for “reference parameters” when using RPC

# (Contrived) example of how call by reference and call by copy-restore can differ

```
int a;  
foo(int x) {  
    x = 2; a = 0;  
}  
int main( ) {  
    foo(a); print(a)  
}
```

Call by reference outputs 0; call  
by copy-restore outputs 2

# Failures

- Many things can go wrong with RPC, e.g., server crash
  - How do we know, from client's perspective, if the server crashed?
  - Supposing we know the server crashed, what do we do from the client side?
    - Run RPC again?
    - Something else?



# Rendezvous

- Similar to RPC
  - Key difference: no new process created on the server
    - Unlike RPC, built-in synchronization between operations
  - Caller side is the same as with RPC
  - Server design is as follows:

in op1(...)

    execute code for op1

[] op2(...)

    execute code for op2

ni

Server blocks until  $\geq 1$  pending invocation on any operation (can be implemented via UNIX *select*)

# Bank Account Problem

- Several people (threads) share a savings account
  - Current balance is sum of all deposits to date minus sum of all withdrawals to date
  - Balance must never become negative
  - Deposits don't delay (except for mutual exclusion)
  - Withdrawals must delay until sufficient funds are in the account

# Bank Account with Monitors

monitor Account

int balance = INIT\_BALANCE

Cond bank

```
Withdraw(amount) {  
    while (amount > balance)  
        Wait(bank)  
    balance = balance - amount  
}
```

```
Deposit(amount) {  
    balance = balance + amount  
    Broadcast(bank)  
}
```

end monitor

# Bank Account with RPC

module Customer // Executes on client

Deposit(amount) or Withdraw(amount) // Customer calls one or the other

module Bank // Executes on server

monitor Account

void Deposit(int)

void Withdraw(int)

int balance = INIT\_BALANCE, Cond bank

Withdraw(int amount)

while (amount > balance)

Wait(bank)

balance = balance – amount

Deposit(int amount)

balance = balance + amount

Broadcast(bank)

Note: Deposit and Withdraw  
need to be monitor functions  
(not shown here; refer to prev. slide)

# Bank Account with Rendezvous

```
module Customer // Executes on client
```

```
    Deposit(amount) or Withdraw(amount) // Remote invocations
```

```
module Bank // Just a class---not a monitor! Executes on server
```

```
    void Deposit(int)
```

```
    void Withdraw(int)
```

```
    int balance = INIT_BALANCE
```

```
    process Teller {
```

```
        while (true)
```

```
            in Deposit(amount)
```

```
                balance = balance + amount
```

```
            [] Withdraw(amount) and balance >= amount
```

```
                balance = balance - amount
```

```
        ni
```

```
    }
```